MONASH University

FIT1043 Introduction to Data Science

Week 5

Ian K T Tan

School of Information Technology
Monash University Malaysia

With materials from Wray Buntine, Mahsa Salehi

Week 4 Coverage
Data Sources
Data Wrangling

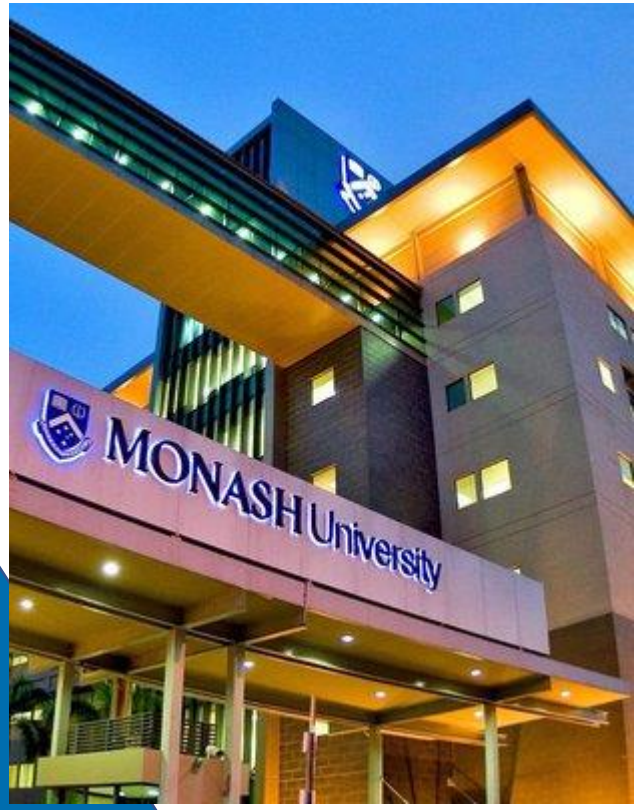

MONASH University

Week 4 Coverage

Data Sources and Wrangling

Open Data

API

Data Quality

Data Auditing

Techniques to handle data quality issues

Week 1

Overview of data science

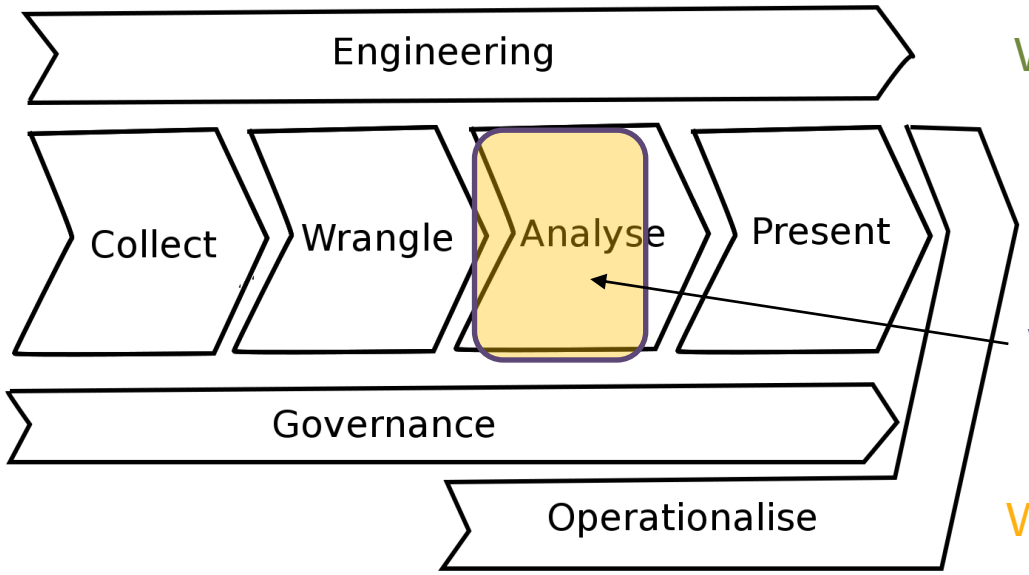

Weeks 9-10

Week 3

Week 4

Weeks 5-7

Week 11

Weeks 2 & 8

Tools for data science

Week	Activities	Assignments
1	Overview of data science	
2	Introduction to Python for data science	
3	Data visualisation and descriptive statistics	
4	Data sources and data wrangling	
5	Data analysis theory	Assignment 1
6	Regression analysis	
7	Classification and clustering	
8	Introduction to R for data science	Assignment 2
9	Characterising data and "big" data	
10	Big data processing	
11	Issues in data management	Assignment 3
12	Industry guest lecture (tentative)	

Week 5 Outline

Introduction to Data Analysis

- What is model?
- What are predictive models?
- How to evaluate predictive models?

Overview of Machine Learning

- Machine learning styles
- What is learning theory
- Linear Regression
- Polynomial regression

Learning Outcomes

Week 5

By the end of this week you should be able to:

- Explain what are models and predictive models
- Analyse predictive models in different examples
- Understand how to evaluate predictive models
- **Analyse how to estimate linear regression model**
- Apply linear regression and polynomial regression on different data sets using Python

Data Model

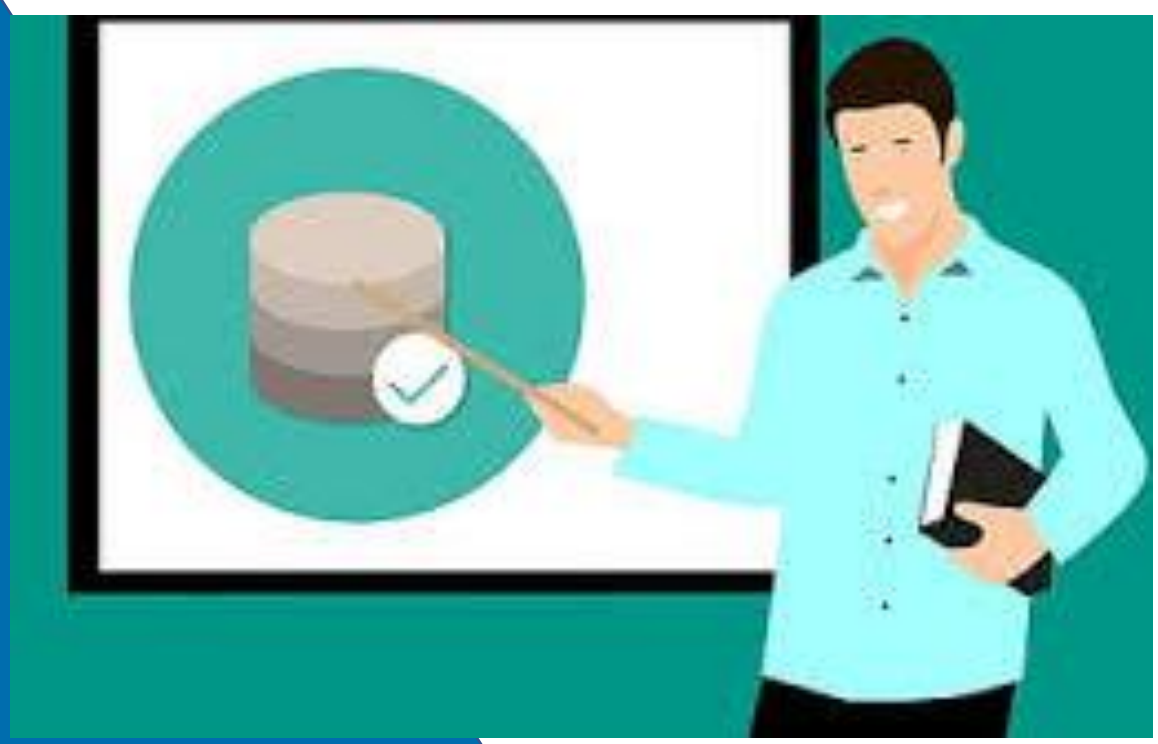

MONASH University

What is a Model?

What is a Model?

Can you draw this …

What is a Model?

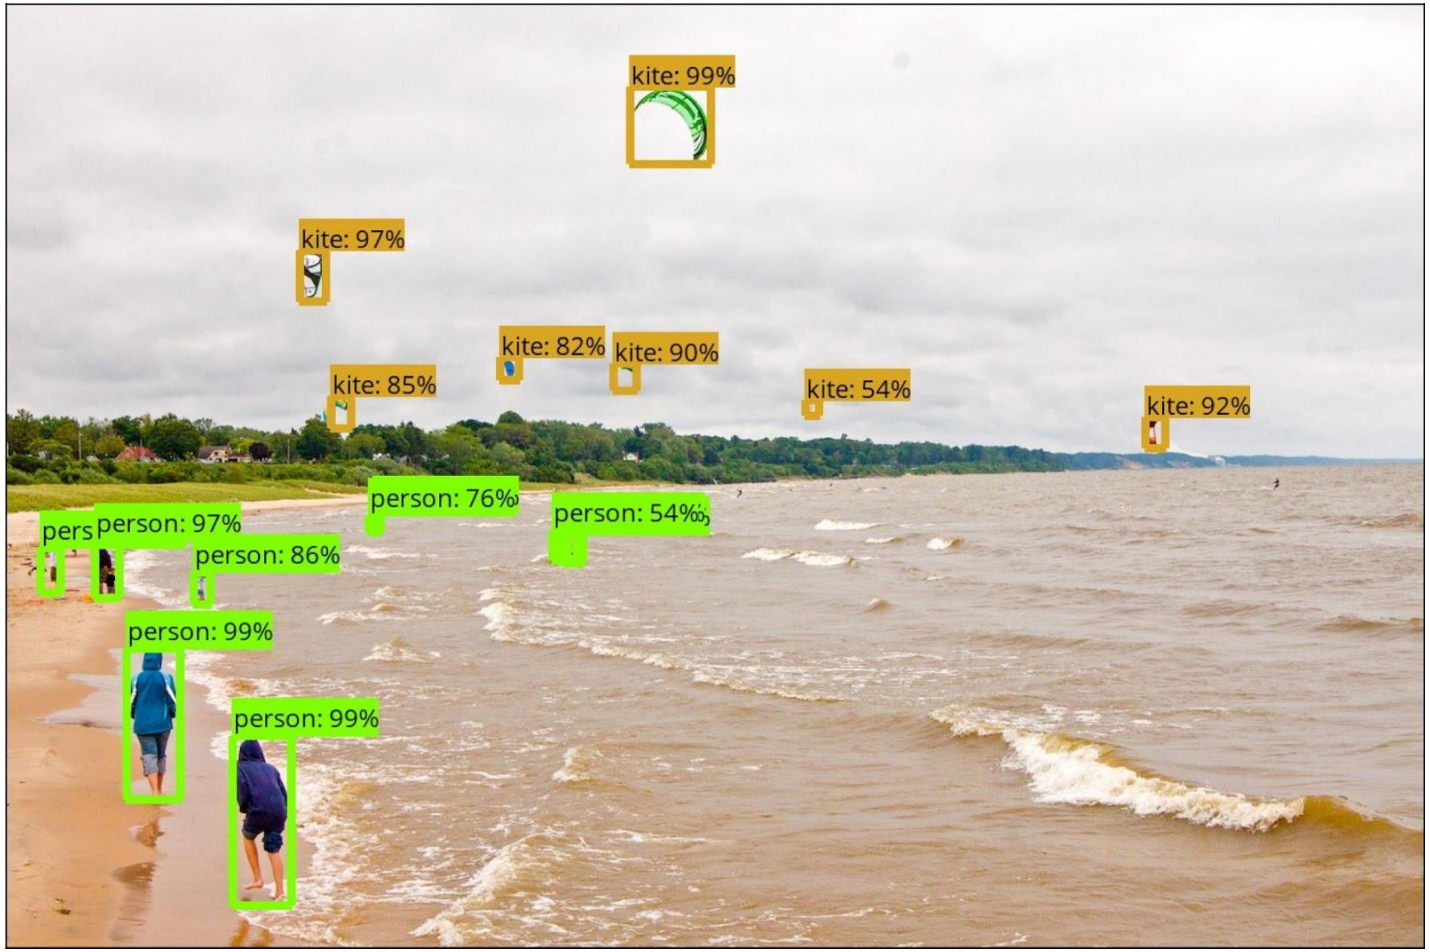

Image Source: https://medium.com/@WuStangDan/step-by-step-tensorflow-object-detection-api-tutorial-part-1-selecting-a-model-a02b6aabe39e

What is a Model?

(1) Help us understand how something works, and
(2) Help us to predict the unknown.

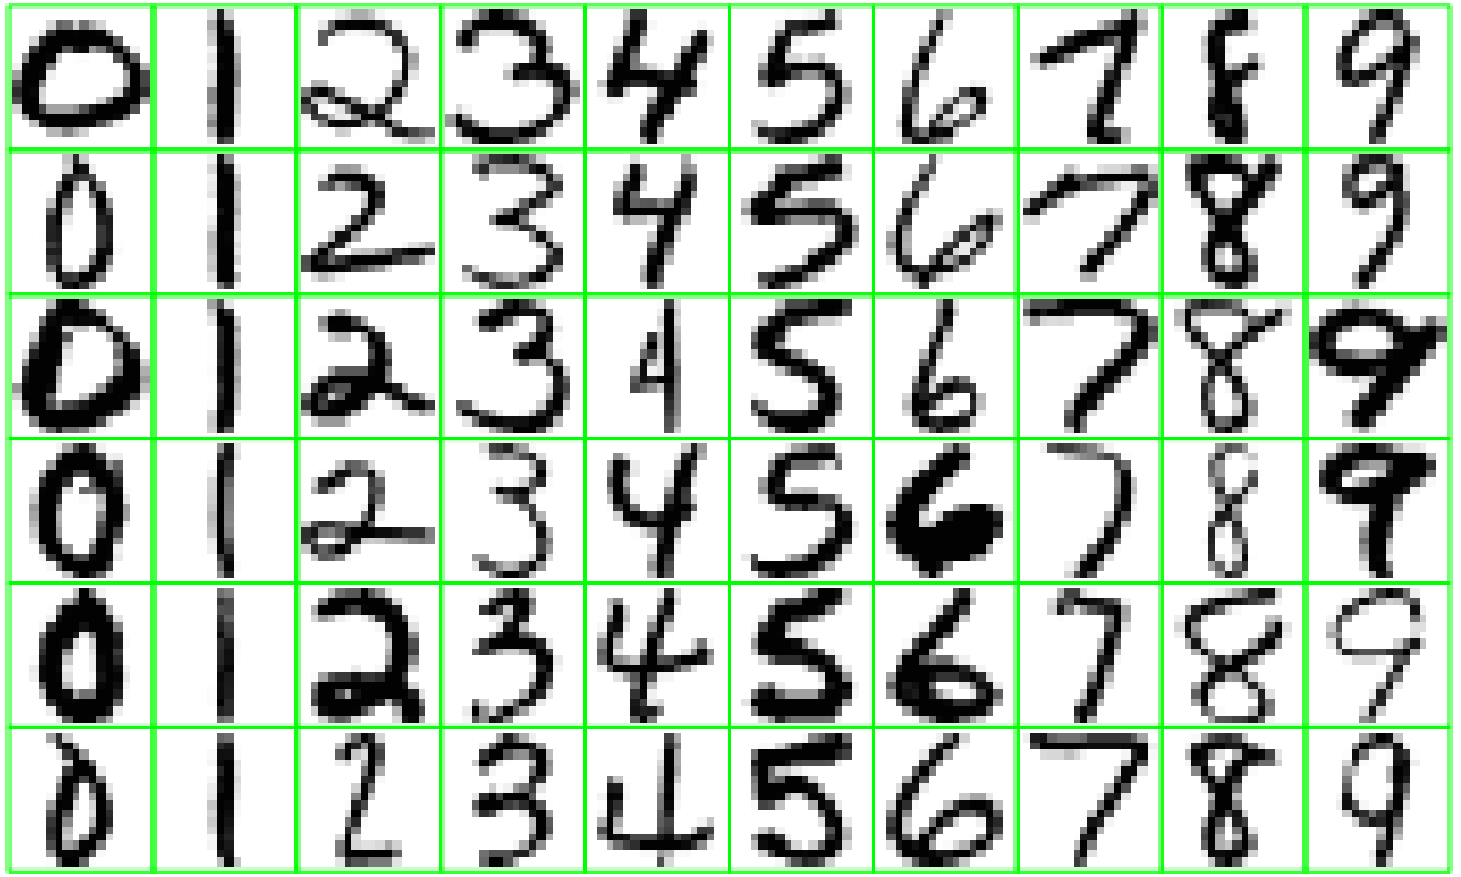

A brief Introduction to Predictive Models For Data Science

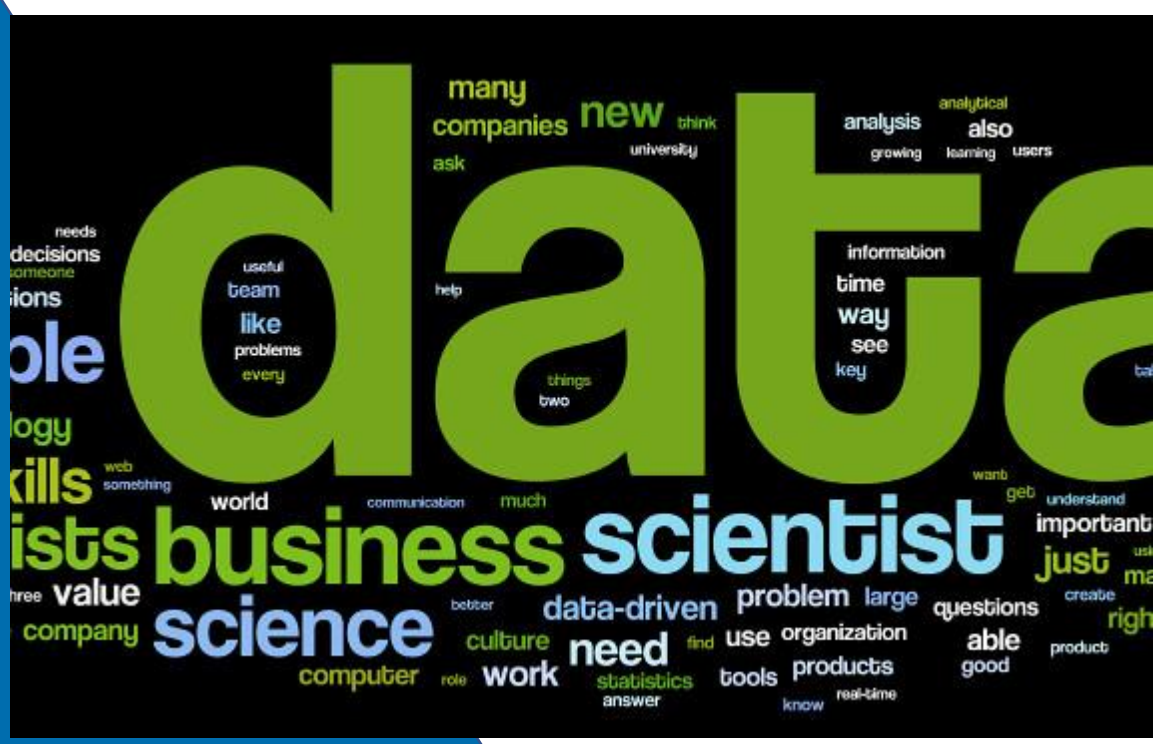

Example from Duda & Hart, Pattern Classification & Scene Analysis, 1973

MONASH University

Motivating Example

On a fishing boat, a conveyor belt loads fish and a worker separates them into their different types

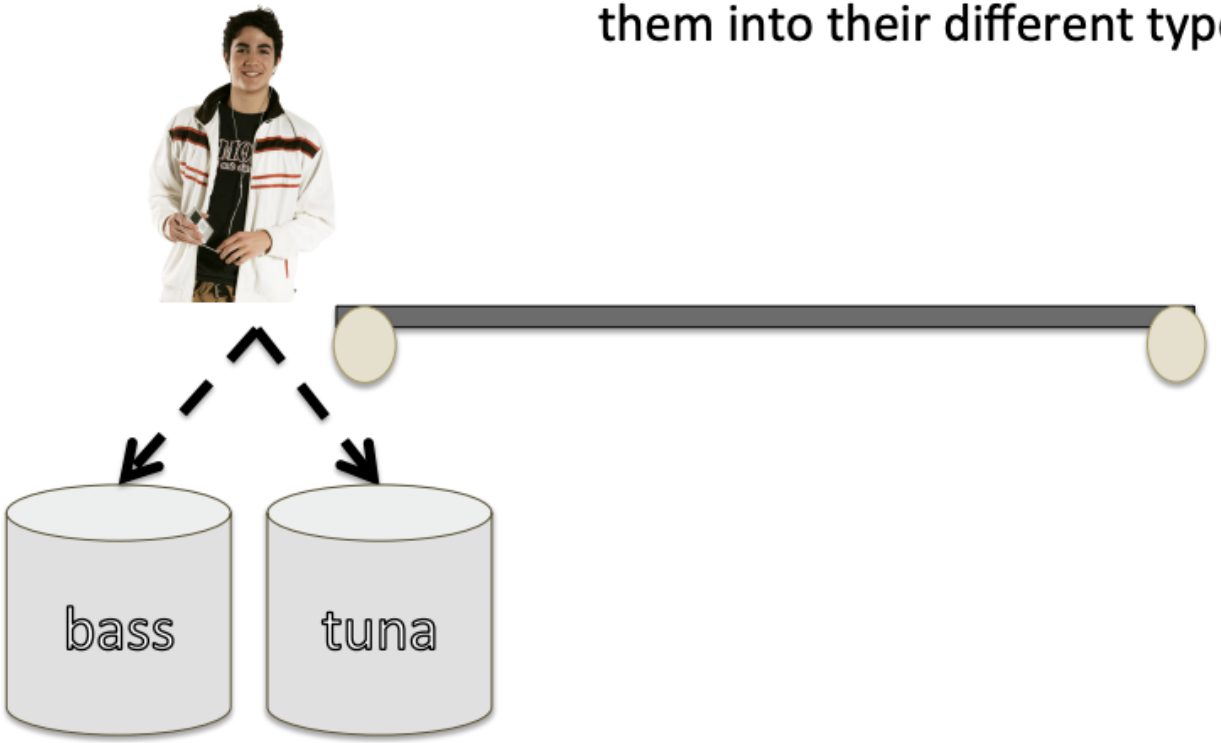

Motivating Example

On a fishing boat, a conveyor belt loads fish and a worker separates them into their different types

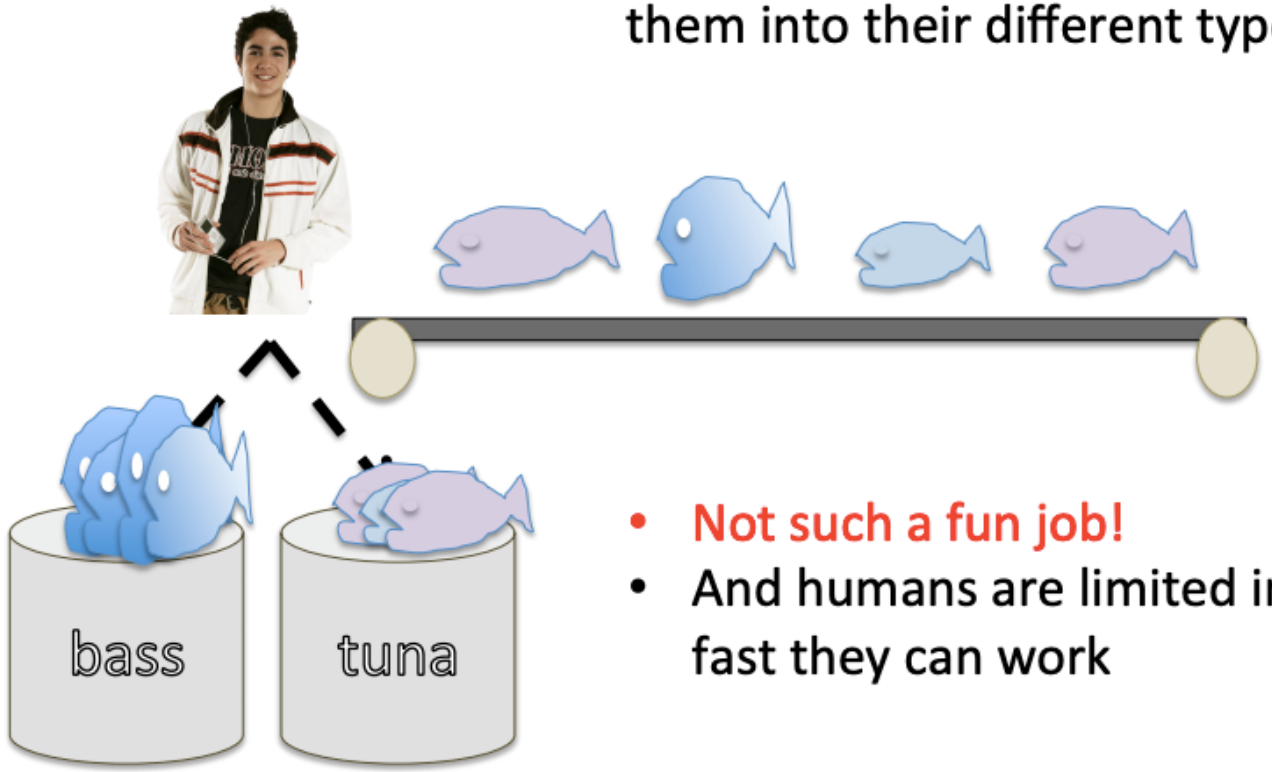

- Not such a fun job!
- And humans are limited in how fast they can work

Motivating Example

Motivating Example

Question: Can we build a system to do the task automatically?

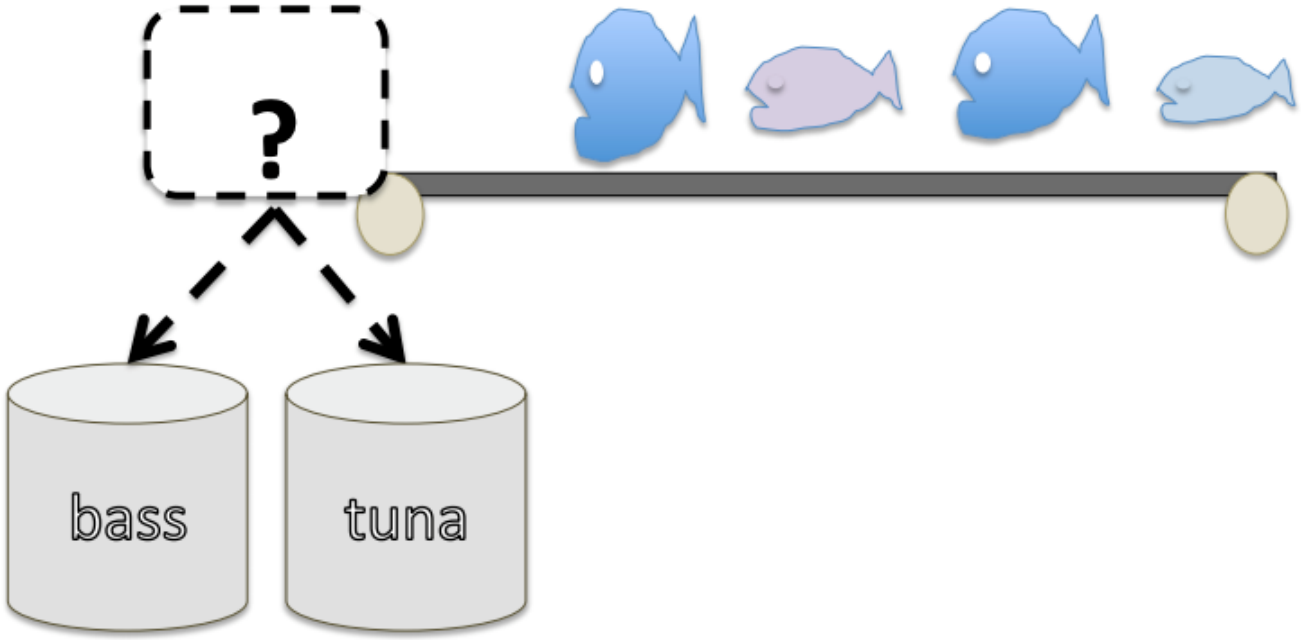

Motivating Example

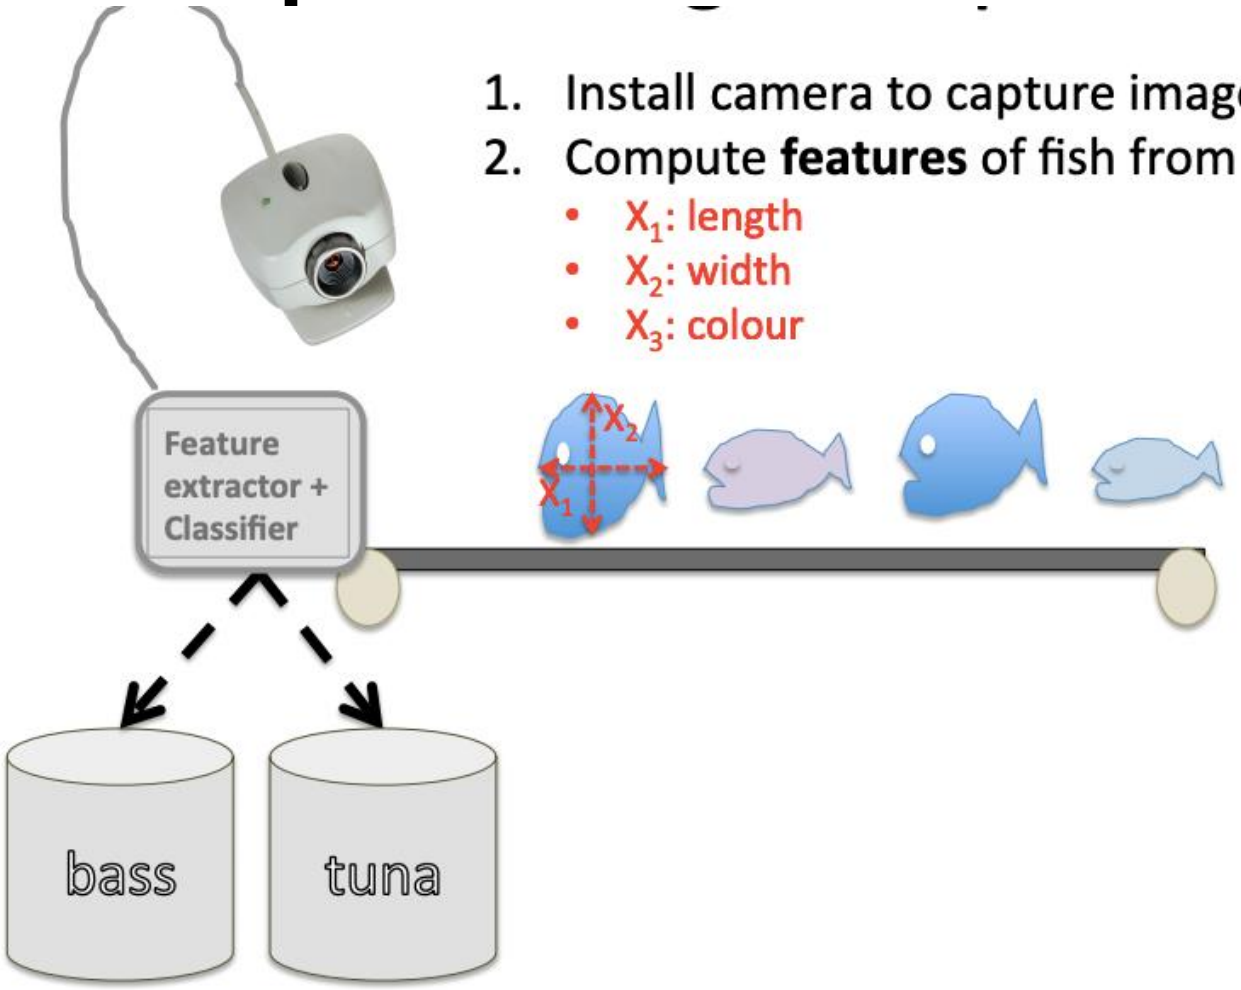

Motivating Example

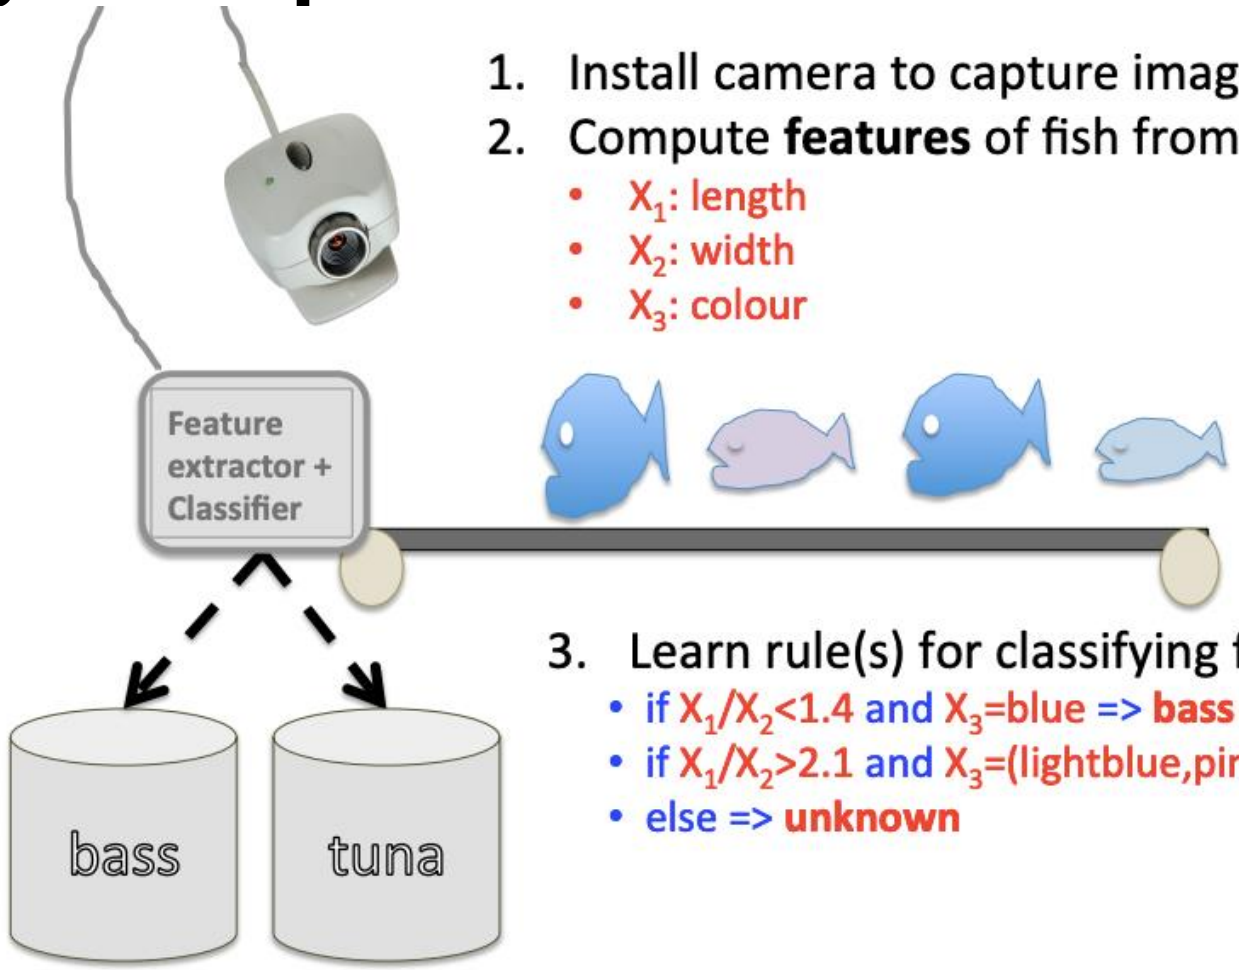

1. Install camera to capture image of fish
2. Compute **features** of fish from image:
   - X_1: length
   - X_2: width
   - X_3: colour

3. Learn rule(s) for classifying fish
   - if $X_1/X_2<1.4$ and X_3=blue => **bass**
   - if $X_1/X_2>2.1$ and X_3=(lightblue,pink) => **tuna**
   - else => **unknown**

Predictive Models

MONASH University

Predictive Models

A predictive model is any model that makes a prediction

- Usually based on a set of features describing an object.
- The prediction could be:
  - A binary outcome (spam, not-spam)
  - Categorical (bass, tuna, other)
  - A real value (the age of the fish)
  - A vector of real values (probability of bass, tuna)
  - Etc.

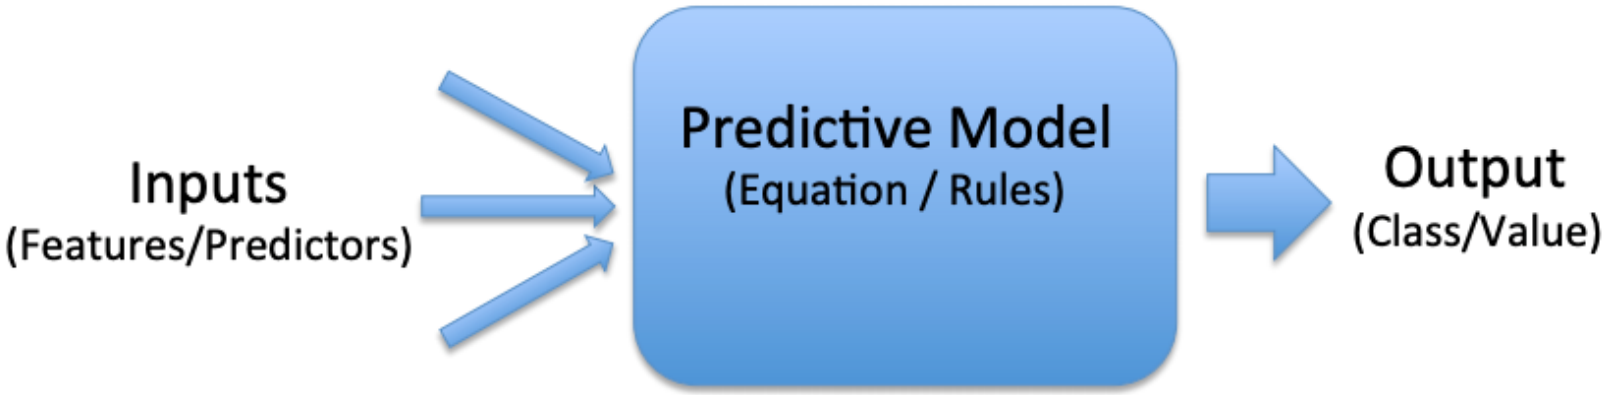

Predictive Models

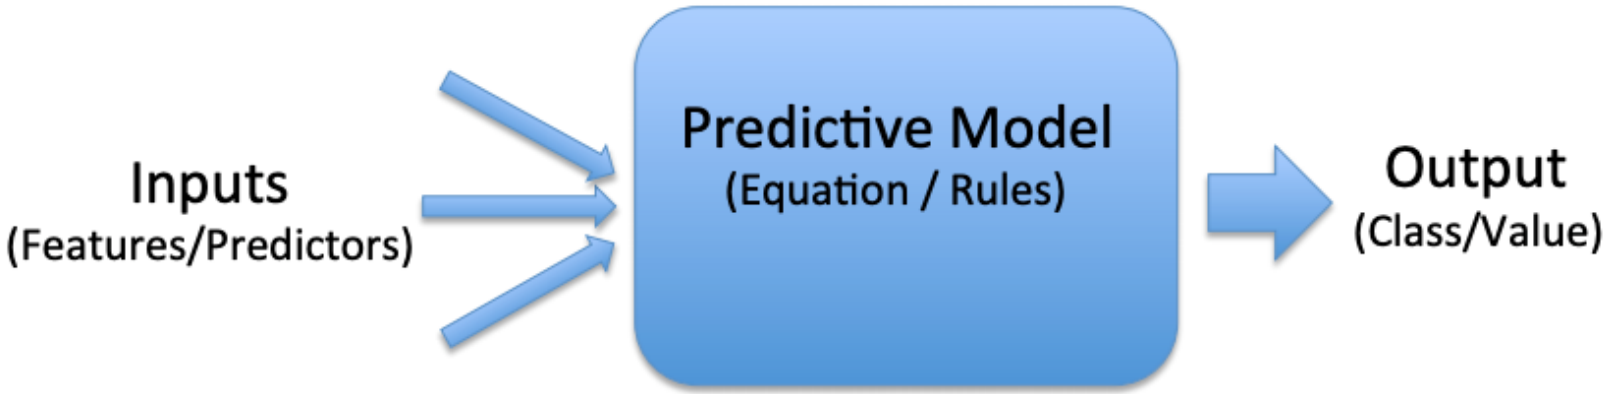

- If the predicted value is binary/categorical we usually refer to the model as a **classifier**
- If it predicts real values we refer to it as **regression**
- Although there are many other types of models, such as rankings, translation (your predictive words) and so on.

Predictive Models

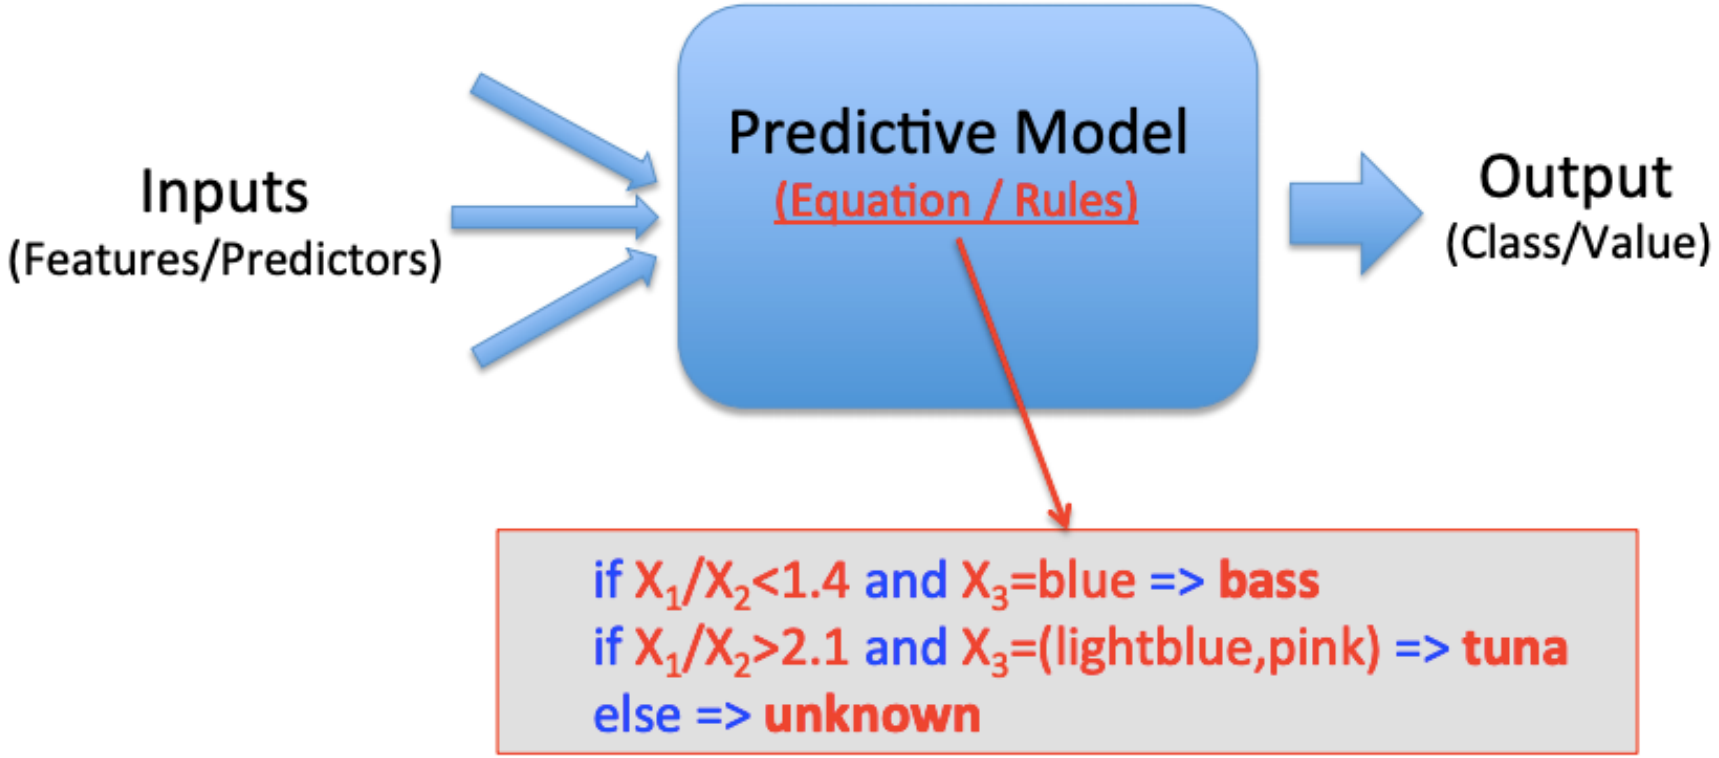

- The predictive model uses **equations**/**rules** to map the input features to output values

Predictive Models

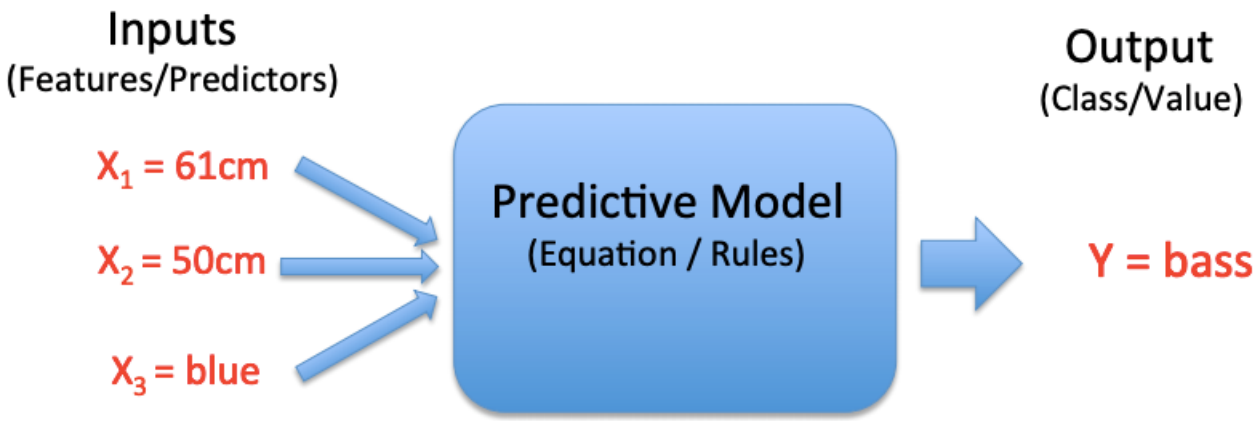

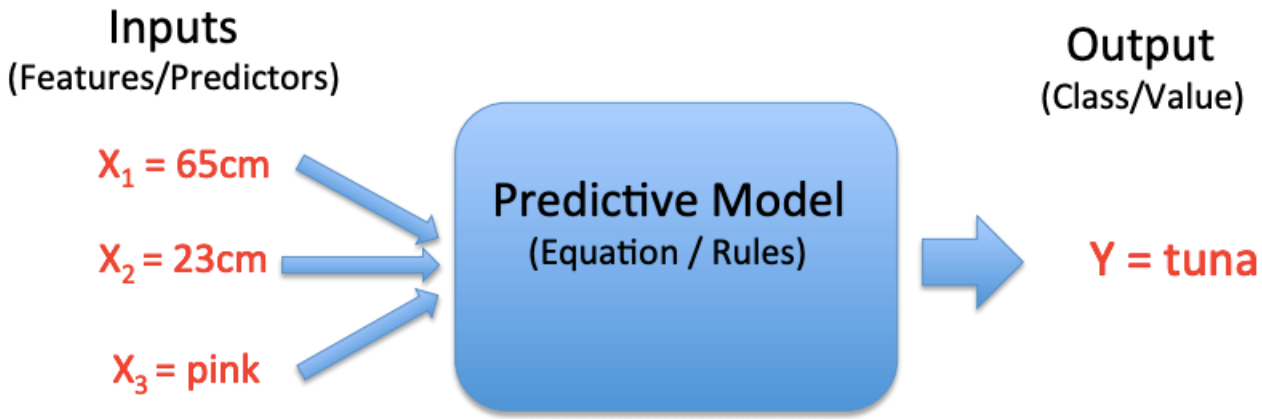

MONASH University

Models are Built from Examples

Most models are developed through learning from examples

Instance	X1 = length	X2 = width	X3 = colour	Y = class
	55	51	blue	bass
	65	23	pink	tuna
	67	54	blue	bass
	54	20	light-blue	tuna
	62	26	pink	tuna
	44	62	blue	bass
	47	55	light-blue	bass
	73	31	pink	tuna
	54	48	light-blue	bass
	57	23	light-blue	tuna

MONASH University

Training and Testing Models

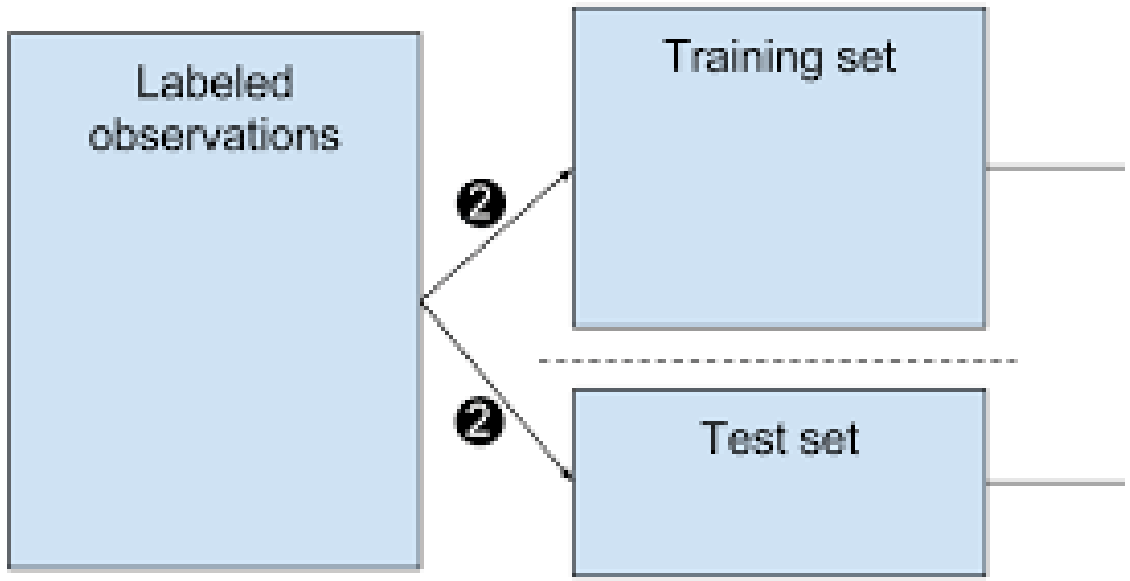

MONASH University

Training a Model

Predictive models are learnt from training data and then applied to make predictions on new instances

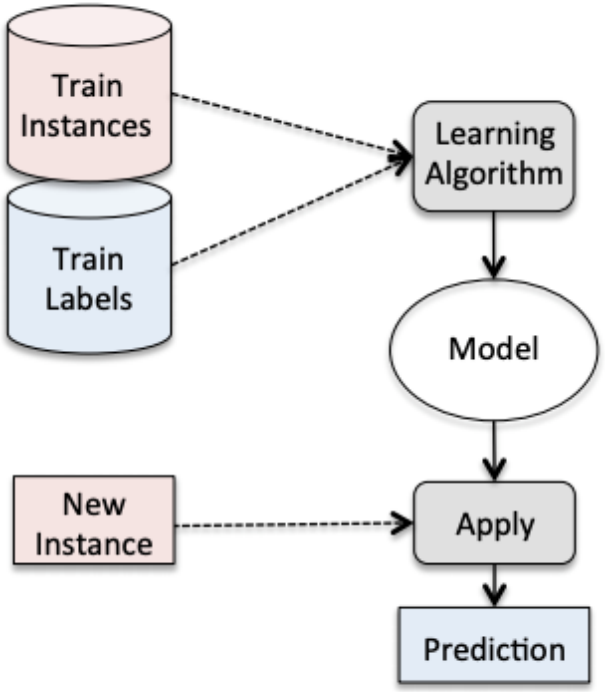

Training a Model

How are models derived (learnt)?

Each training instance (fish in our case) is just a point in some feature space

Here the colour denotes the class

- blue = bass,
- green = tuna,
- red = unknown

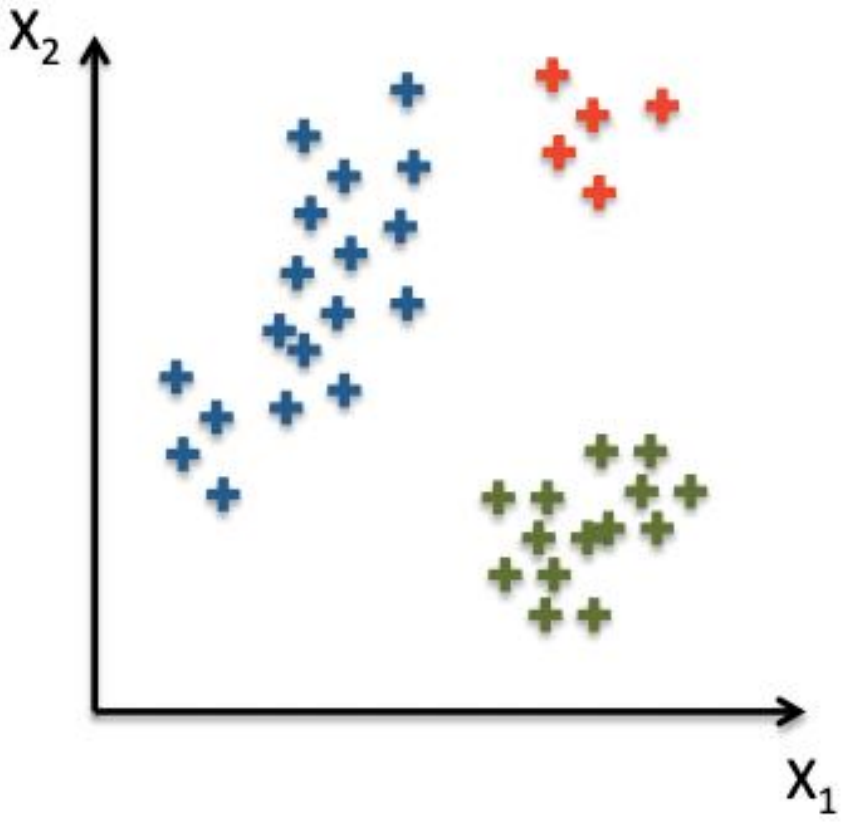

Training a Model

How are models derived (learnt)?

Many (classification) learning algorithms work by **dividing the feature space into regions of the same type**

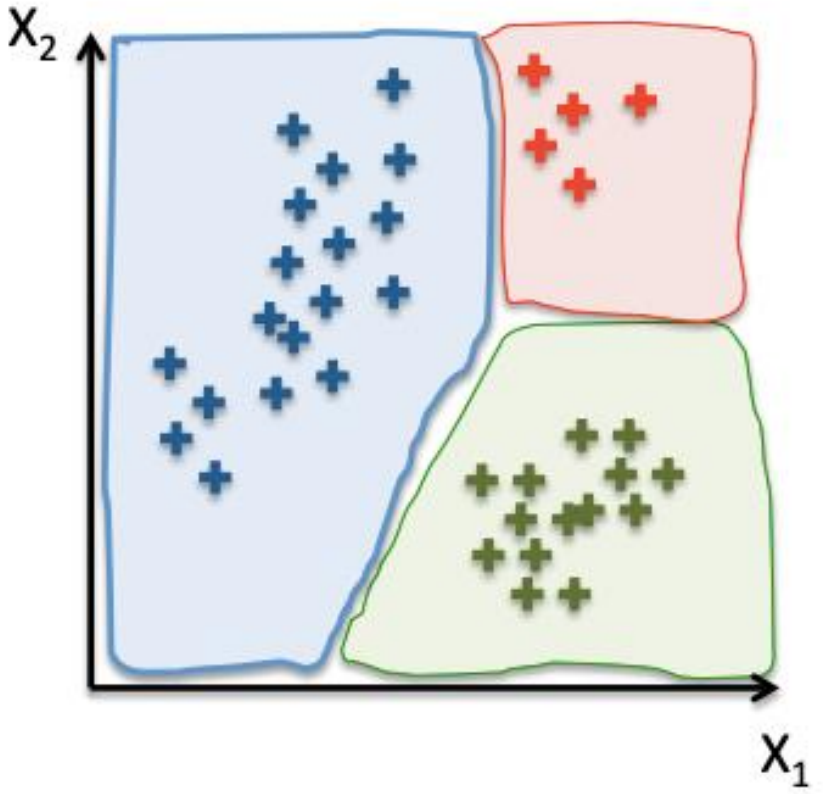

In Practise

In practice, the data is usually **overlapping**

Making it hard to separate the classes

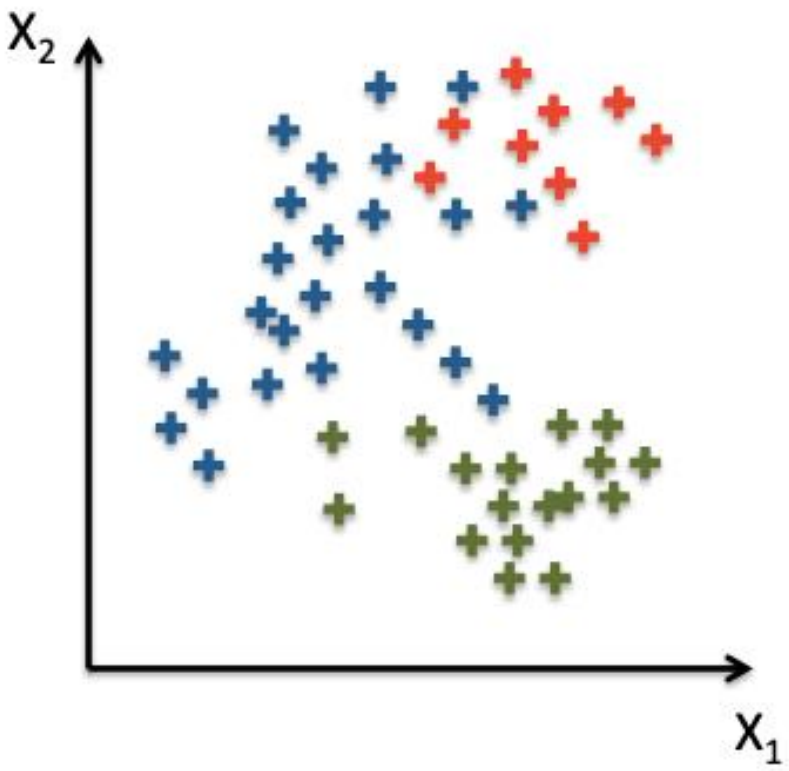

In Practise

And we have many feature dimensions

With **some features more useful than others**

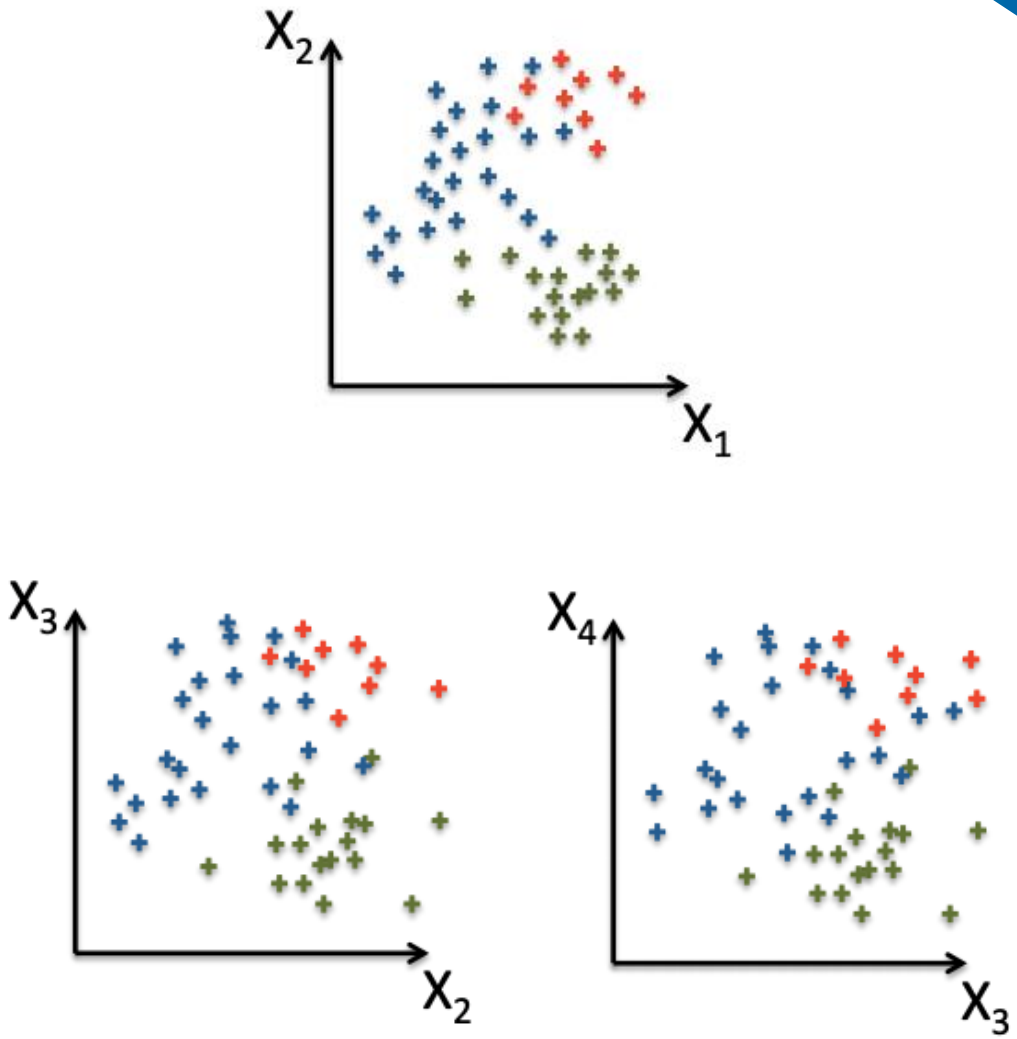

Different Models

There are many different types of models that we can train to classify objects

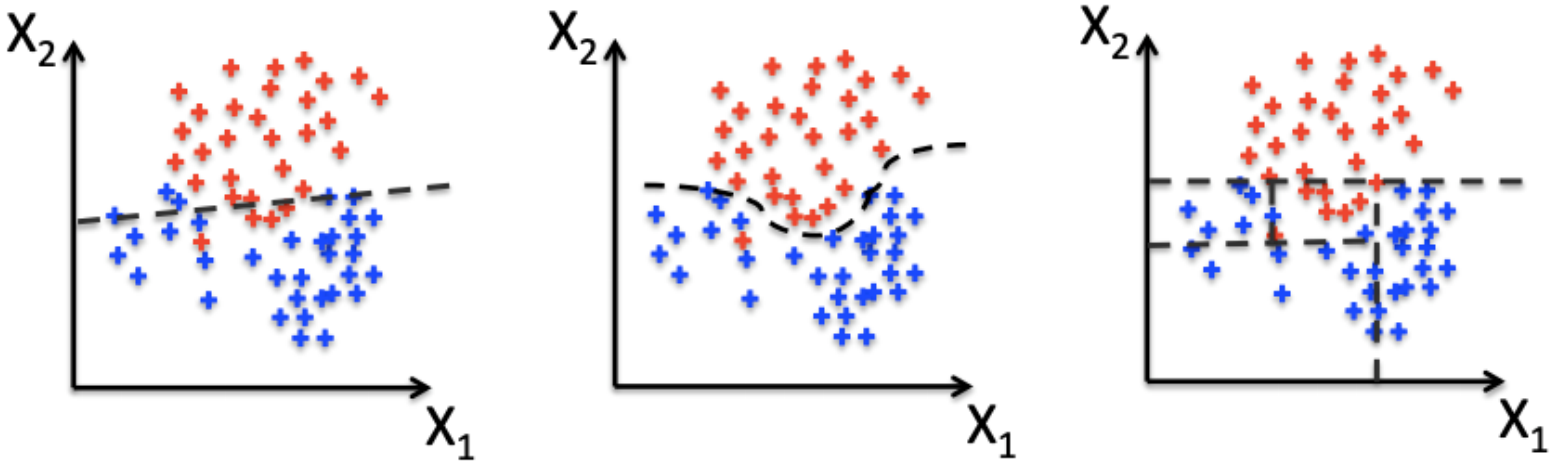

Linear classifiers

e.g. Logistic Regression, Linear SVMs

Non-linear Classifiers

e.g. Neural Nets, SVM with RBF kernel

Decision Tree Learners

e.g. Random forests

Different Models

How can we decide which model is better?

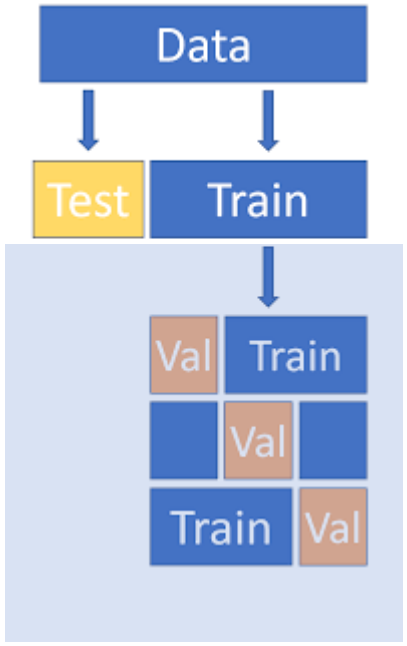

Testing Models

We evaluate predictive models based on how well they predict the labels for test instances (not used in training)

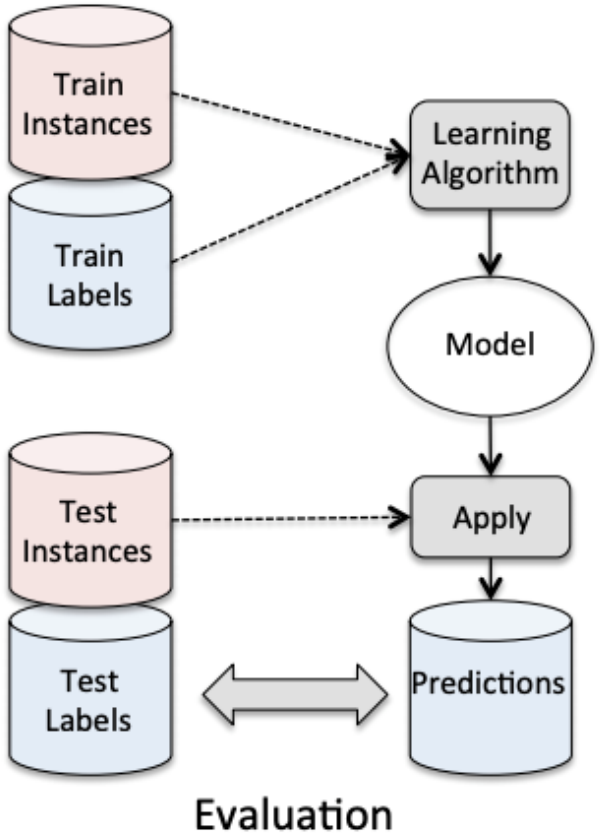

Performance of Predictive Models

Generally:

- The **more training data the better** the test performance
- And (providing there is sufficient training data) the more features the better performance will be
  - Caveat: There is a limit to how many features

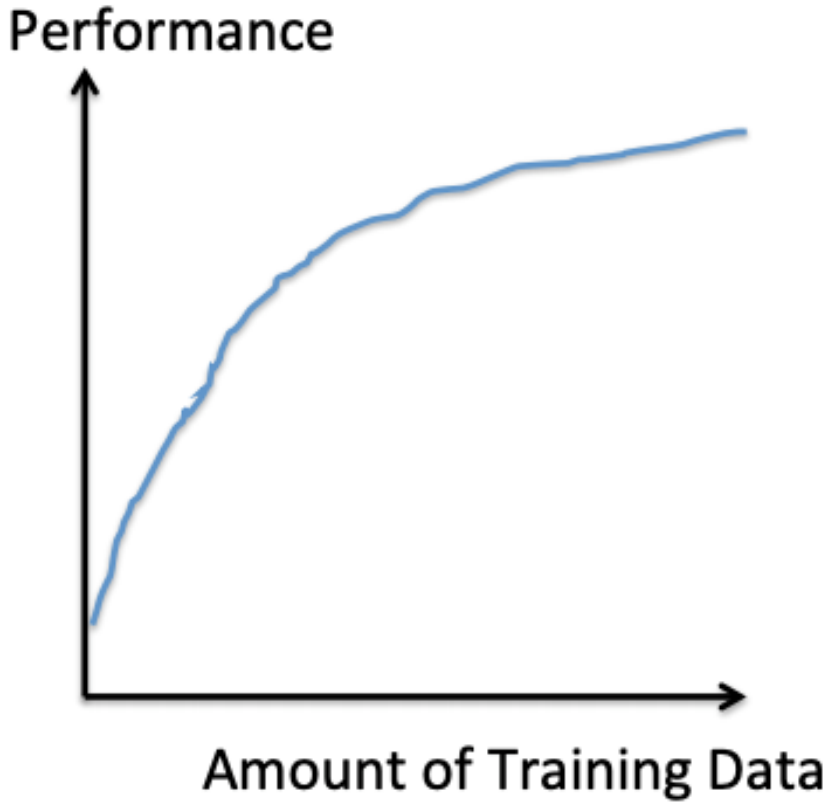

Home Activities

Suggested Activities for the week

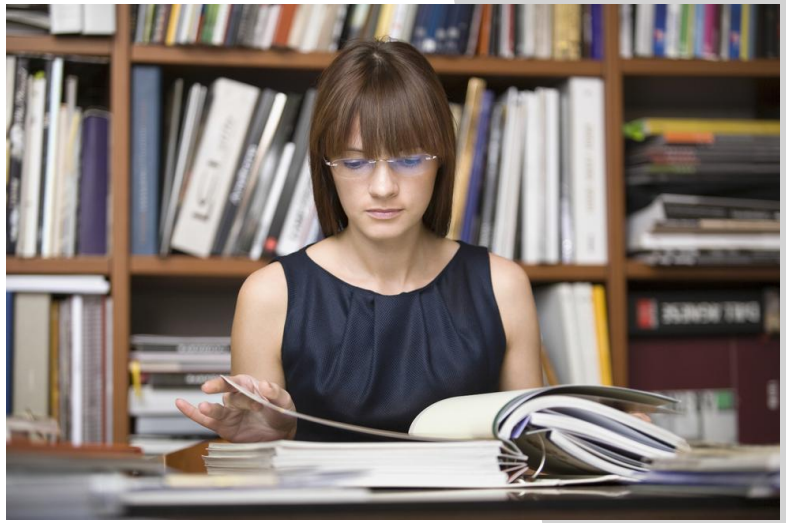

Videos

Watch Jeremy Howard, “The wonderful and terrifying implications of computers that can learn”, TEDxBrussels

Watch Fei-Fei Li, “How we are teaching computers to understand pictures”, TED2015

Articles

Read Tarang Shah, “About Train, Validation and Test Sets in Machine Learning”, towards data science

MONASH University

Recap: Learning Outcomes

Week

By the end of this week you should be able to:

- Explain what are models and predictive models
- Analyse predictive models in different examples
- Understand how to evaluate predictive models
- **Analyse how to estimate linear regression model**
- Apply linear regression and polynomial regression on different data sets using Python